

TECTEAM2 / LATAM / LIVE DEMO



PRISM

PAYMENT INCIDENT INTELLIGENCE

**Detect, diagnose and quantify
conversion drops before they cost money.**

SIGNAL IN. EVIDENCE OUT.

Conversion drops in silence.

A single drop hides inside thousands of transactions — across countries, providers, methods and issuing banks.

01 **Manual investigation**

Teams pivot across dashboards to find where it broke.

02 **Slow diagnosis**

Root cause takes hours while exposure compounds.

03 **Reactive response**

Action arrives after the revenue is already lost.

THE FIVE QUESTIONS

WHAT dropped

WHERE it dropped

WHY it dropped

WHO is affected

HOW MUCH is at risk

Every minute spent investigating increases financial exposure.

One symptom in. Evidence out.

Continuous multi-segment traffic, evaluated end to end in under two minutes.

01	02	03	04	05
Detect	Validate	Diagnose	Quantify	Recommend
Adaptive 1 & 5-minute windows, FDR-corrected statistics.	Filters normal fluctuation before it reaches anyone.	Root cause by provider, issuer and decline code.	Impact normalised to USD, plus projected cost.	Per-team actions, human approval, audit trail.

PIPELINE

scan > validate > cluster > root_cause > impact > priority > human_review

AI accelerates investigation. Humans remain in control of the decision.

Seven layers, one evidence trail.

01	Data	60-day history plus live multi-segment traffic
02	Baselines & aggregation	Adaptive 1- and 5-minute windows per detection level
03	Detection	Statistical scan with FDR correction and temporal persistence
04	Diagnosis	Root cause by provider, merchant, payment method and decline code
05	Impact	Financial quantification normalised to USD, projected cost
06	Prioritisation	Explainable P1–P4 score with operational guardrails
07	Interface & control	FastAPI, dashboard, AI narrative, human approval with audit trail

What does an unresolved payment incident cost?

PRISM translates a conversion drop into current exposure, business impact, and the cost of delayed resolution.

01 — MEASURE THE GAP

Expected vs. actual approvals

PRISM compares the incident's live approval rate with its historical baseline, then estimates the approvals that were lost.

02 — QUANTIFY TODAY'S EXPOSURE

Value at risk, adjusted for recovery

Lost approvals are valued using the average transaction amount in USD. Expected retry recovery is deducted to estimate the exposure that remains unrecovered.

03 — PROJECT THE COST OF DELAY

4 h, 24 h, and 7 days

The current exposure rate is projected across three horizons. A 6-hour MTTR assumption distinguishes the full downside of an unresolved incident from the expected impact if the team restores service on time.

4 HOURS

Incident assumed active across the full horizon.

24 HOURS

Exposure keeps accruing. Expected net impact assumes resolution after 6 hours.

7 DAYS

Shows the longer-term downside. Expected net impact remains limited by the 6-hour MTTR assumption.

The model reports three business views: GMV at risk, platform revenue at risk, and merchant economic impact.

6 filters before anyone is paged.

01	≥ 30 attempts	Ignores segments with too little traffic to judge.
02	Approval drop ≥ 5pp	Requires a materially worse approval rate.
03	Reliable baseline	Needs enough history to compare against.
04	Z-score ≥ 2	Demands statistical significance, not noise.
05	FDR Benjamini-Hochberg	Adjusts p-values to control the False Discovery Rate (FDR).
05	Persistence ≥ 2 windows	Must hold across consecutive windows before alerting.

When PRISM will not claim a cause.

- | | | |
|----|-----------------------------------|---|
| 01 | Confidence score < 0.70 | Evidence is not strong enough to confirm the diagnosis. |
| 02 | Unresolved candidate | The anomaly is stored as a candidate, not as a confirmed incident. |
| 03 | Human review | Operations reviews the evidence and decides whether to investigate further. |
| 04 | Continue monitoring | The system keeps watching new transactions until confidence improves or the candidate is dismissed. |

No evidence, no claim — the incident stays a candidate until it earns confidence.

Which incident goes first.

An explainable score, not an opinion — and not simply the largest anomaly.

SCORE WEIGHTS

40%

Financial impact

20%

Confidence

15%

Duration

10%

Merchant scope

10%

Root cause scope

5%

Merchant criticality

PRIORITY BANDS

P1

Critical

Immediate attention

P2

High

High-priority investigation

P3

Medium

Investigate & monitor

P4

Low

Continue monitoring

Three decisions, and their cost.

#01

T+03:53

When does a conversion drop become an incident?

→ Historical baseline + adaptive thresholds — deviation, confidence, volume, persistence.

Fixed thresholds ignore normal differences across merchants and providers; pure statistics flags events that are financially irrelevant. We accepted slower reaction in low-volume segments to cut alert fatigue.

#02

T+10:54

Automatic remediation or human approval?

→ Human approval — only approved actions execute, and every decision is logged.

Full automation would minimise time to mitigation, but a wrong diagnosis could worsen conversion. We accepted a small delay in exchange for accountability and a complete audit trail.

#03

T+11:01

What happens to low-confidence anomalies?

→ Surfaced as “under investigation” — promoted only at confidence ≥ 0.70 .

Forcing a root cause risks sending a team to the wrong component; dropping candidates silently hides real signals. PRISM can state that the evidence is not yet sufficient.

PRISM / PAYMENT INCIDENT INTELLIGENCE

PRISM

SIGNAL CAUTION2

CRITICAL SIGNAL1

LIVE SIMULATION

06:55:05

ANALYST VIEW

OP

LIVE NOW

REPLAY: BEFORE THE INCIDENTS (10:00-10:02)

TRIAL BY FIRE

3 ACTIVE INCIDENTS

Multiple live incidents detected and prioritized

Control Tower separated the active failures, estimated their financial impact, and generated operator recommendations.

UPDATED 10:11 • 6-HOUR BASELINE

NETWORK CONVERSION ⓘ

89.4%

-8.6 pp

Expected90.0%

ACTIVE INCIDENTS

3

Separated

Highest priorityP1

GMV AT RISK

\$73.0K

USD / hour

Confidence adjusted\$9.4K

MONITORED TRAFFIC

14.4K

100%

Merchants / providers3 / 3

EXECUTIVE SUMMARY

Network-wide financial exposure

ACROSS ALL ACTIVE INCIDENTS

TOTAL ECONOMIC IMPACT

\$2.3K

Active incidents3

GMV AT RISK (ADJUSTED)

\$9.4K

Gross, unadjusted\$14.4K

PLATFORM REVENUE AT RISK

\$66

Fee revenue exposure-

MERCHANT ECONOMIC IMPACT

\$2.2K

Merchants' lost margin-

WHERE THE CURRENT GMV AT RISK CONCENTRATES - DIAGNOSTIC SIGNAL, NOT THE CONFIRMED ROOT CAUSE

ROOT CAUSE

provider 75%

decline_code 17%

issuing_bank 7%

COUNTRY

BR 93%

MX 7%

PROVIDER

Adyen 75%

ISSUING BANK

BBVA 7%

DECLINE CODE

PROCESSOR_ERROR 93%

ISSUER_UNAVAILABLE 7%

ASK PRISM

Automated evidence, human decision.,

01 Continuous monitoring

Control Tower watches performance continuously.

02 Statistical detection

Detects significant anomalies and validates that they persist.


03 Root cause & impact

Diagnoses the affected segment and estimates financial impact.

04 Evidence-based recommendation

Prioritises and presents the action to a human operator.

Detection and diagnosis are automated; the decision stays with the operator.



TECTEAM2 / LATAM / LIVE DEMO



PRISM

PAYMENT INCIDENT INTELLIGENCE

**Detect, diagnose and quantify
conversion drops before they cost money.**

SIGNAL IN. EVIDENCE OUT.